

Assignment 1 – Part 1 : Classification using K-Nearest Neighbors (KNN)

The goal of this task is to perform a classification problem using Telescope dataset with the K-Nearest Neighbors (KNN) algorithm.

We will: Balance the dataset. Split it into training, validation, and testing sets (70%, 15%, 15%). Implement KNN manually (from scratch).

Re-implement using Scikit-Learn. Tune the parameter k and plot validation accuracy. Compare both implementations. Evaluate with accuracy, precision, recall, F1-score, and confusion matrix. Discuss overfitting and underfitting behavior.

Importing Libraries

```
import csv
import pandas as pd
from sklearn.model_selection import train_test_split
import numpy as np
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

Load and Explore the Dataset

```
df=pd.read_csv("telescope_data.csv")

print(df.head())
print(f"Rows= {df.shape[0]}")
print(f"Columns= {df.shape[1]}")
```

```
   Unnamed: 0  fLength  fWidth  fSize  fConc  fConc1  fAsym  fM3Long  \
0            0    28.7967   16.0021  2.6449  0.3918  0.1982   27.7004  22.0110
1            1    31.6036   11.7235  2.5185  0.5303  0.3773   26.2722  23.8238
2            2   162.0520  136.0310  4.0612  0.0374  0.0187  116.7410 -64.8580
3            3    23.8172    9.5728  2.3385  0.6147  0.3922   27.2107  -6.4633
4            4    75.1362   30.9205  3.1611  0.3168  0.1832   -5.5277  28.5525

   fM3Trans  fAlpha  fDist  class
0   -8.2027  40.0920   81.8828    g
1   -9.9574   6.3609  205.2610    g
2  -45.2160  76.9600  256.7880    g
3   -7.1513  10.4490  116.7370    g
4   21.8393   4.6480  356.4620    g
Rows= 19020
Columns= 12
```

Data Cleaning

We remove: Duplicate rows Missing values Unused columns (like Unnamed: 0) to ensure data consistency.

```
df.drop_duplicates()
df.dropna(axis=0)
#axis=1 columns axis=0 rows
```

	Unnamed: 0	fLength	fWidth	fSize	fConc	fConc1	fAsym	fM3Long	fM3Trans	fAlpha	fDist	class
0	0	28.7967	16.0021	2.6449	0.3918	0.1982	27.7004	22.0110	-8.2027	40.0920	81.8828	g
1	1	31.6036	11.7235	2.5185	0.5303	0.3773	26.2722	23.8238	-9.9574	6.3609	205.2610	g
2	2	162.0520	136.0310	4.0612	0.0374	0.0187	116.7410	-64.8580	-45.2160	76.9600	256.7880	g
3	3	23.8172	9.5728	2.3385	0.6147	0.3922	27.2107	-6.4633	-7.1513	10.4490	116.7370	g
4	4	75.1362	30.9205	3.1611	0.3168	0.1832	-5.5277	28.5525	21.8393	4.6480	356.4620	g
...	...	...	...	...	...	...	...	...	...	...	...	...
19015	19015	21.3846	10.9170	2.6161	0.5857	0.3934	15.2618	11.5245	2.8766	2.4229	106.8258	h
19016	19016	28.9452	6.7020	2.2672	0.5351	0.2784	37.0816	13.1853	-2.9632	86.7975	247.4560	h
19017	19017	75.4455	47.5305	3.4483	0.1417	0.0549	-9.3561	41.0562	-9.4662	30.2987	256.5166	h
19018	19018	120.5135	76.9018	3.9939	0.0944	0.0683	5.8043	-93.5224	-63.8389	84.6874	408.3166	h
19019	19019	187.1814	53.0014	3.2093	0.2876	0.1539	-167.3125	-168.4558	31.4755	52.7310	272.3174	h

19020 rows x 12 columns

```
df.drop(columns=['Unnamed: 0'], inplace=True)
df.head()
print(df.shape)
```

(19020, 11)

Handle Class Imbalance

The dataset contains more gamma (g) than hadron (h) events. We will undersample the gamma class to match the hadron count.

```
gamma=df[df['class']=='g']
hadron=df[df['class']=='h']
print(gamma.shape)
print(hadron.shape)
```

(12332, 11)
(6688, 11)

```
gamma_sample=gamma.sample(hadron.shape[0],random_state=42)
print(gamma_sample.shape)
```

(6688, 11)

```
balanced=pd.concat([gamma_sample,hadron])
balanced=balanced.sample(frac=1,random_state=42).reset_index(drop=True)
#frac=1 -> all 100%
#drop=true drop unused rows
```

```
print(len(balanced[balanced['class']=='g']))
```

```
6688
```

Data Splitting

We split the dataset into: 70% training 15% validation 15% testing Stratified sampling ensures class balance in each subset.

```
X=balanced.drop(columns=['class'])
Y=balanced['class']
x_train,x_temp,y_train,y_temp=train_test_split(X,Y,test_size=0.3,random_state=42,stratify=Y)
print(x_train.shape)
print(y_train.shape)
```

```
(9363, 10)
(9363,)
```

```
x_validate,x_test,y_validate,y_test=train_test_split(x_temp,y_temp,test_size=0.5,random_state=42,stratify=y_temp)
print(x_test.shape)
print(y_test.shape)
print(x_validate.shape)
print(y_validate.shape)
```

```
(2007, 10)
(2007,)
(2006, 10)
(2006,)
```

Manual KNN Implementation

We now implement KNN manually: Compute Euclidean distance Select the K-nearest neighbors Use majority voting for classification

```
def eclidean_distance(x1,x2):
    return np.sqrt(np.sum((x1-x2)**2))
```

```
class KNN:
    def __init__(self, k=3):
        self.k = k

    def fit(self, X, y):
        self.X_train = X.to_numpy()
        self.y_train = y.to_numpy()

    def predict(self, X):
        predictions = []
        for x in X.to_numpy():
            distances = [eclidean_distance(x, x_train) for x_train in self.X_train]
```

```

        k_indices = np.argsort(distances)[:self.k]
        k_nearest_labels = self.y_train[k_indices]
        # Majority vote
        label = max(set(k_nearest_labels), key=list(k_nearest_labels).count)
        predictions.append(label)
    return np.array(predictions)

```

```

knn = KNN(k=5)
knn.fit(x_train, y_train)

y_pred = knn.predict(x_validate)

accuracy = np.mean(y_pred == y_validate.to_numpy())
print(f"Validation Accuracy: {accuracy:.4f}")

```

Validation Accuracy: 0.7692

Validate Manual KNN and Tune k

We test multiple odd k-values (up to $\sqrt{n}$) to find the best k.

```

n=len(x_train)
k_max=int(np.sqrt(n))
k_values=range(1,k_max+1,2)
accuracies = []

for k in k_values:
    knn = KNN(k=k)
    knn.fit(x_train, y_train)
    y_pred = knn.predict(x_validate)
    acc = np.mean(y_pred == y_validate.to_numpy())
    accuracies.append(acc)
    print(f"k = {k}, Validation Accuracy = {acc:.4f}")

# Best k
best_k = k_values[np.argmax(accuracies)]
print(f"\n Best k = {best_k} with accuracy = {max(accuracies):.4f}")

```

```

k = 1, Validation Accuracy = 0.7453
k = 3, Validation Accuracy = 0.7557
k = 5, Validation Accuracy = 0.7692
k = 7, Validation Accuracy = 0.7742
k = 9, Validation Accuracy = 0.7762
k = 11, Validation Accuracy = 0.7737
k = 13, Validation Accuracy = 0.7687
k = 15, Validation Accuracy = 0.7697
k = 17, Validation Accuracy = 0.7717
k = 19, Validation Accuracy = 0.7732
k = 21, Validation Accuracy = 0.7732
k = 23, Validation Accuracy = 0.7712
k = 25, Validation Accuracy = 0.7687
k = 27, Validation Accuracy = 0.7717
k = 29, Validation Accuracy = 0.7757
k = 31, Validation Accuracy = 0.7767


```

```
k = 33, Validation Accuracy = 0.7762
k = 35, Validation Accuracy = 0.7752
k = 37, Validation Accuracy = 0.7812
k = 39, Validation Accuracy = 0.7742
k = 41, Validation Accuracy = 0.7782
k = 43, Validation Accuracy = 0.7802
k = 45, Validation Accuracy = 0.7772
k = 47, Validation Accuracy = 0.7742
k = 49, Validation Accuracy = 0.7792
k = 51, Validation Accuracy = 0.7722
k = 53, Validation Accuracy = 0.7757
k = 55, Validation Accuracy = 0.7722
k = 57, Validation Accuracy = 0.7752
k = 59, Validation Accuracy = 0.7767
k = 61, Validation Accuracy = 0.7747
k = 63, Validation Accuracy = 0.7762
k = 65, Validation Accuracy = 0.7772
k = 67, Validation Accuracy = 0.7747
k = 69, Validation Accuracy = 0.7717
k = 71, Validation Accuracy = 0.7712
k = 73, Validation Accuracy = 0.7707
k = 75, Validation Accuracy = 0.7697
k = 77, Validation Accuracy = 0.7682
k = 79, Validation Accuracy = 0.7667
k = 81, Validation Accuracy = 0.7687
k = 83, Validation Accuracy = 0.7667
k = 85, Validation Accuracy = 0.7667
k = 87, Validation Accuracy = 0.7662
k = 89, Validation Accuracy = 0.7677
k = 91, Validation Accuracy = 0.7672
k = 93, Validation Accuracy = 0.7642
k = 95, Validation Accuracy = 0.7682
```

Best k = 37 with accuracy = 0.7812

Plot Validation Accuracy (Manual KNN)

```
plt.plot(k_values, accuracies, marker='o')
plt.xlabel('k')
plt.ylabel('Validation Accuracy')
plt.title('K (odd) vs. Accuracy')
plt.grid(True)
plt.show()
```



Z-Normalization

```
x_z_norm= (X - X.mean(axis=0)) / X.std(axis=0)
x_z_norm_train,x_z_norm_temp,y_z_norm_train,y_z_norm_temp=train_test_split(x_z_norm,Y,test_size=0.3,random_state=42,stratify=Y)
x_z_norm_validate,x_z_norm_test,y_z_norm_validate,y_z_norm_test=train_test_split(x_z_norm_temp,y_z_norm_temp,test_size=0.5,random_state=42,stratify=y_z_norm_temp)
```

Visualizing how z-Normalization Change the data Before normalization, some features have very different ranges, which can bias distance-based algorithms like KNN. After applying Z-score normalization, all features are rescaled around mean = 0 and standard deviation = 1.



```
fig, ax = plt.subplots(1, 2, figsize=(12,4))

# Before normalization
x_train.iloc[:,0].hist(ax=ax[0], bins=30, color='red')
ax[0].set_title('Feature Example Before Normalization')
ax[0].set_xlabel('Feature Value')
ax[0].set_ylabel('Frequency')

# After normalization
x_z_norm_train.iloc[:,0].hist(ax=ax[1], bins=30, color='green')
ax[1].set_title('Feature Example After Normalization')
ax[1].set_xlabel('Feature Value (Normalized)')
ax[1].set_ylabel('Frequency')

plt.suptitle('Effect of Normalization on Feature Distribution', fontsize=13)
plt.tight_layout()
plt.show()
```

Effect of Normalization on Feature Distribution



Manual KNN with Normalized Data

We repeat KNN on normalized features for fairer distance calculation.

```
n=len(x_z_norm_train)
k_max=int(np.sqrt(n))
k_values=range(1,k_max+1,2)
accuracies_z_norm = []

for k in k_values:
    knn = KNN(k=k)
    knn.fit(x_z_norm_train, y_z_norm_train)
    y_z_norm_pred = knn.predict(x_z_norm_validate)
    acc = np.mean(y_z_norm_pred == y_z_norm_validate.to_numpy())
    accuracies_z_norm.append(acc)
    print(f"k = {k}, Validation Accuracy = {acc:.4f}")

# Best k
best_k_z_norm = k_values[np.argmax(accuracies_z_norm)]
print(f"\nBest k = {best_k_z_norm} with accuracy = {max(accuracies_z_norm):.4f}")
```


```
k = 1, Validation Accuracy = 0.7971
k = 3, Validation Accuracy = 0.8121
k = 5, Validation Accuracy = 0.8175
k = 7, Validation Accuracy = 0.8180
k = 9, Validation Accuracy = 0.8161
k = 11, Validation Accuracy = 0.8225
k = 13, Validation Accuracy = 0.8156
k = 15, Validation Accuracy = 0.8156
k = 17, Validation Accuracy = 0.8166
k = 19, Validation Accuracy = 0.8136
k = 21, Validation Accuracy = 0.8146
k = 23, Validation Accuracy = 0.8121
k = 25, Validation Accuracy = 0.8111
k = 27, Validation Accuracy = 0.8106
```

```
k = 29, Validation Accuracy = 0.8131
k = 31, Validation Accuracy = 0.8166
k = 33, Validation Accuracy = 0.8126
k = 35, Validation Accuracy = 0.8131
k = 37, Validation Accuracy = 0.8116
k = 39, Validation Accuracy = 0.8111
k = 41, Validation Accuracy = 0.8096
k = 43, Validation Accuracy = 0.8116
k = 45, Validation Accuracy = 0.8081
k = 47, Validation Accuracy = 0.8046
k = 49, Validation Accuracy = 0.8031
k = 51, Validation Accuracy = 0.8051
k = 53, Validation Accuracy = 0.8046
k = 55, Validation Accuracy = 0.8031
k = 57, Validation Accuracy = 0.8036
k = 59, Validation Accuracy = 0.8026
k = 61, Validation Accuracy = 0.8026
k = 63, Validation Accuracy = 0.8051
k = 65, Validation Accuracy = 0.8046
k = 67, Validation Accuracy = 0.8031
k = 69, Validation Accuracy = 0.8016
k = 71, Validation Accuracy = 0.8021
k = 73, Validation Accuracy = 0.8021
k = 75, Validation Accuracy = 0.8006
k = 77, Validation Accuracy = 0.8016
k = 79, Validation Accuracy = 0.8026
k = 81, Validation Accuracy = 0.8036
k = 83, Validation Accuracy = 0.8041
k = 85, Validation Accuracy = 0.8041
k = 87, Validation Accuracy = 0.8021
k = 89, Validation Accuracy = 0.8021
k = 91, Validation Accuracy = 0.8016
k = 93, Validation Accuracy = 0.8026
k = 95, Validation Accuracy = 0.8021
```

Best k = 11 with accuracy = 0.8225

Plot Validation Accuracy (Manual with Normalization)

```
plt.plot(k_values, accuracies_z_norm, marker='o')
plt.xlabel('k')
plt.ylabel('Validation Accuracy')
plt.title('K (odd) vs. Accuracy using Z normalization')
plt.grid(True)
plt.show()
```



Apply KNN using Scikit-Learn

We now use KNeighborsClassifier for comparison.

```
knn = KNeighborsClassifier(n_neighbors=11)
knn.fit(x_z_norm_train, y_z_norm_train)
y_pred = knn.predict(x_z_norm_validate)
acc = accuracy_score(y_z_norm_validate, y_pred)
print(f"Accuracy: {acc:.4f}")
```

Accuracy: 0.8225

It generates same accuracy as our implementation

Sklearn model with no normalization

```
n=len(x_z_norm_train)
k_max=int(np.sqrt(n))
k_values=range(1,k_max+1,2)
accuracies_sklearn_non_norm = []
for k in k_values:
    # Create KNN model
    knn = KNeighborsClassifier(n_neighbors=k)

    # Train on training data
    knn.fit(x_train, y_train)

    # Validate on validation set
```

```
y_val_pred = knn.predict(x_validate)

# Compute accuracy
acc = accuracy_score(y_validate, y_val_pred)
accuracies_sklearn_non_norm.append(acc)
print(f"k = {k}, Validation Accuracy = {acc:.4f}")


# Find best k
best_k_sklearn_non_norm = k_values[np.argmax(accuracies_sklearn_non_norm)]
print(f"\nBest k = {best_k_sklearn_non_norm} with accuracy = {max(accuracies_sklearn_non_norm):.4f}")
```

```
k = 1, Validation Accuracy = 0.7453
k = 3, Validation Accuracy = 0.7557
k = 5, Validation Accuracy = 0.7692
k = 7, Validation Accuracy = 0.7742
k = 9, Validation Accuracy = 0.7762
k = 11, Validation Accuracy = 0.7737
k = 13, Validation Accuracy = 0.7687
k = 15, Validation Accuracy = 0.7697
k = 17, Validation Accuracy = 0.7717
k = 19, Validation Accuracy = 0.7732
k = 21, Validation Accuracy = 0.7732
k = 23, Validation Accuracy = 0.7712
k = 25, Validation Accuracy = 0.7687
k = 27, Validation Accuracy = 0.7717
k = 29, Validation Accuracy = 0.7757
k = 31, Validation Accuracy = 0.7767
k = 33, Validation Accuracy = 0.7762
k = 35, Validation Accuracy = 0.7752
k = 37, Validation Accuracy = 0.7812
k = 39, Validation Accuracy = 0.7742
k = 41, Validation Accuracy = 0.7782
k = 43, Validation Accuracy = 0.7802
k = 45, Validation Accuracy = 0.7772
k = 47, Validation Accuracy = 0.7742
k = 49, Validation Accuracy = 0.7792
k = 51, Validation Accuracy = 0.7722
k = 53, Validation Accuracy = 0.7757
k = 55, Validation Accuracy = 0.7722
k = 57, Validation Accuracy = 0.7752
k = 59, Validation Accuracy = 0.7767
k = 61, Validation Accuracy = 0.7747
k = 63, Validation Accuracy = 0.7762
k = 65, Validation Accuracy = 0.7772
k = 67, Validation Accuracy = 0.7747
k = 69, Validation Accuracy = 0.7717
k = 71, Validation Accuracy = 0.7712
k = 73, Validation Accuracy = 0.7707
k = 75, Validation Accuracy = 0.7697
k = 77, Validation Accuracy = 0.7682
k = 79, Validation Accuracy = 0.7667
k = 81, Validation Accuracy = 0.7687
k = 83, Validation Accuracy = 0.7667
k = 85, Validation Accuracy = 0.7667
k = 87, Validation Accuracy = 0.7662
k = 89, Validation Accuracy = 0.7677
k = 91, Validation Accuracy = 0.7672
k = 93, Validation Accuracy = 0.7642
k = 95, Validation Accuracy = 0.7682
```

```
Best k = 37 with accuracy = 0.7812
```

Plot Validation Accuracy (Sklern KNN without Normalization)

```
plt.plot(k_values, accuracies_sklern_non_norm, marker='o')
plt.xlabel('k')
plt.ylabel('Accuracy')
plt.title('K Value vs Accuracy using sklern KNN with no normalization')
plt.grid(True)
plt.show()
```



Sklearn model with Z-normalization

```
accuracies_sklern = []
for k in k_values:
    # Create KNN model
    knn = KNeighborsClassifier(n_neighbors=k)

    # Train on training data
    knn.fit(x_z_norm_train, y_z_norm_train)

    # Validate on validation set
    y_val_pred = knn.predict(x_z_norm_validate)

    # Compute accuracy
    acc = accuracy_score(y_z_norm_validate, y_val_pred)
    accuracies_sklern.append(acc)
    print(f"k = {k}, Validation Accuracy = {acc:.4f}")

# Find best k
```

```
best_k_sklearn = k_values[np.argmax(accuracies_sklearn)]
print(f"\nBest k = {best_k_sklearn} with accuracy = {max(accuracies_sklearn):.4f}")
```


```
k = 1, Validation Accuracy = 0.7971
k = 3, Validation Accuracy = 0.8121
k = 5, Validation Accuracy = 0.8175
k = 7, Validation Accuracy = 0.8180
k = 9, Validation Accuracy = 0.8161
k = 11, Validation Accuracy = 0.8225
k = 13, Validation Accuracy = 0.8156
k = 15, Validation Accuracy = 0.8156
k = 17, Validation Accuracy = 0.8166
k = 19, Validation Accuracy = 0.8136
k = 21, Validation Accuracy = 0.8146
k = 23, Validation Accuracy = 0.8121
k = 25, Validation Accuracy = 0.8111
k = 27, Validation Accuracy = 0.8106
k = 29, Validation Accuracy = 0.8131
k = 31, Validation Accuracy = 0.8166
k = 33, Validation Accuracy = 0.8126
k = 35, Validation Accuracy = 0.8131
k = 37, Validation Accuracy = 0.8116
k = 39, Validation Accuracy = 0.8111
k = 41, Validation Accuracy = 0.8096
k = 43, Validation Accuracy = 0.8116
k = 45, Validation Accuracy = 0.8081
k = 47, Validation Accuracy = 0.8046
k = 49, Validation Accuracy = 0.8031
k = 51, Validation Accuracy = 0.8051
k = 53, Validation Accuracy = 0.8046
k = 55, Validation Accuracy = 0.8031
k = 57, Validation Accuracy = 0.8036
k = 59, Validation Accuracy = 0.8026
k = 61, Validation Accuracy = 0.8026
k = 63, Validation Accuracy = 0.8051
k = 65, Validation Accuracy = 0.8046
k = 67, Validation Accuracy = 0.8031
k = 69, Validation Accuracy = 0.8016
k = 71, Validation Accuracy = 0.8021
k = 73, Validation Accuracy = 0.8021
k = 75, Validation Accuracy = 0.8006
k = 77, Validation Accuracy = 0.8016
k = 79, Validation Accuracy = 0.8026
k = 81, Validation Accuracy = 0.8036
k = 83, Validation Accuracy = 0.8041
k = 85, Validation Accuracy = 0.8041
k = 87, Validation Accuracy = 0.8021
k = 89, Validation Accuracy = 0.8021
k = 91, Validation Accuracy = 0.8016
k = 93, Validation Accuracy = 0.8026
k = 95, Validation Accuracy = 0.8021
```

Best k = 11 with accuracy = 0.8225

Plot Validation Accuracy (Sklearn KNN with Normalization)

```
plt.plot(k_values, accuracies_sklearn, marker='o')
plt.xlabel('k')
plt.ylabel('Accuracy')
plt.title('K Value vs Accuracy using sklearn KNN with Normalization')
```

```
plt.grid(True)
plt.show()
```




Effect of Normalization on accuracy

```
# --- Plotting to Compare Normalization Effect ---
plt.figure(figsize=(10, 6))

# Plot accuracy without normalization
plt.plot(k_values, accuracies_sklearn_non_norm, marker='x', linestyle='--', label='Without Normalization', color='red')

# Plot accuracy with normalization
plt.plot(k_values, accuracies_sklearn, marker='o', label='With Z-Score Normalization', color='blue')

plt.title('Effect of Normalization on K-NN Validation Accuracy')
plt.xlabel('k Value')
plt.ylabel('Validation Accuracy')
plt.legend()
plt.grid(True)
plt.show()
```



Trying another distance calculation function as manhattan and check accuracies Manhattan distance = $\sum(\text{abs}(x_1 - x_2))$

```
n=len(x_z_norm_train)
k_max=int(np.sqrt(n))
k_values=range(1,k_max+1,2)
accuracies_sklearn_manhattan= []
for k in k_values:
    # Create KNN model
    knn = KNeighborsClassifier(n_neighbors=k,metric="manhattan")
    #sum(abs(x1 - x2))

    # Train on training data
    knn.fit(x_train, y_train)

    # Validate on validation set
    y_val_pred = knn.predict(x_validate)

    # Compute accuracy
    acc = accuracy_score(y_validate, y_val_pred)
    accuracies_sklearn_manhattan.append(acc)
    print(f"k = {k}, Validation Accuracy = {acc:.4f}")


# Find best k
best_k_sklearn_manhattan = k_values[np.argmax(accuracies_sklearn_manhattan)]
print(f"\nBest k = {best_k_sklearn_manhattan} with accuracy = {max(accuracies_sklearn_manhattan):.4f}")
```

```
k = 1, Validation Accuracy = 0.7498
k = 3, Validation Accuracy = 0.7642
k = 5, Validation Accuracy = 0.7747
k = 7, Validation Accuracy = 0.7807
k = 9, Validation Accuracy = 0.7861
k = 11, Validation Accuracy = 0.7876
k = 13, Validation Accuracy = 0.7846
k = 15, Validation Accuracy = 0.7832
k = 17, Validation Accuracy = 0.7856
k = 19, Validation Accuracy = 0.7861
k = 21, Validation Accuracy = 0.7866
k = 23, Validation Accuracy = 0.7846
k = 25, Validation Accuracy = 0.7861
k = 27, Validation Accuracy = 0.7856
k = 29, Validation Accuracy = 0.7871
k = 31, Validation Accuracy = 0.7841
k = 33, Validation Accuracy = 0.7871
k = 35, Validation Accuracy = 0.7866
k = 37, Validation Accuracy = 0.7861
k = 39, Validation Accuracy = 0.7861
k = 41, Validation Accuracy = 0.7866
k = 43, Validation Accuracy = 0.7876
k = 45, Validation Accuracy = 0.7856
k = 47, Validation Accuracy = 0.7851
k = 49, Validation Accuracy = 0.7817
k = 51, Validation Accuracy = 0.7777
k = 53, Validation Accuracy = 0.7802
k = 55, Validation Accuracy = 0.7792
k = 57, Validation Accuracy = 0.7802
k = 59, Validation Accuracy = 0.7782
k = 61, Validation Accuracy = 0.7752
k = 63, Validation Accuracy = 0.7797
k = 65, Validation Accuracy = 0.7752
k = 67, Validation Accuracy = 0.7747
k = 69, Validation Accuracy = 0.7772
k = 71, Validation Accuracy = 0.7782
k = 73, Validation Accuracy = 0.7792
k = 75, Validation Accuracy = 0.7832
k = 77, Validation Accuracy = 0.7827
k = 79, Validation Accuracy = 0.7836
k = 81, Validation Accuracy = 0.7802
k = 83, Validation Accuracy = 0.7812
k = 85, Validation Accuracy = 0.7827
k = 87, Validation Accuracy = 0.7817
k = 89, Validation Accuracy = 0.7797
k = 91, Validation Accuracy = 0.7792
k = 93, Validation Accuracy = 0.7812
k = 95, Validation Accuracy = 0.7812
```

Best k = 11 with accuracy = 0.7876

Plotting Accuracies

```
plt.plot(k_values, accuracies_sklearn_manhattan, marker='o')
plt.xlabel('k')
plt.ylabel('Accuracy')
plt.title('K Value vs Accuracy using sklearn KNN with Manhattan Distance')
plt.grid(True)
plt.show()
```




Comparing between Validation accuracies using euclidean distance and manhattan distance

```
plt.figure(figsize=(10, 6))

plt.plot(k_values, accuracies_sklearn, marker='o', label='Euclidean Distance (L2)')
plt.plot(k_values, accuracies_sklearn_manhattan, marker='s', linestyle='--', label='Manhattan Distance (L1)')


plt.title('K-NN Validation Accuracy: Euclidean vs. Manhattan Distance')
plt.xlabel('k Value')
plt.ylabel('Validation Accuracy')
plt.legend()
plt.grid(True)
plt.show()
```



Euclidean distance has higher accuracy for the same k

Plotting validation accuracy for both implementations

```
plt.figure(figsize=(8,5))
plt.plot(k_values, accuracies_z_norm, marker='o', label='Manual KNN (Z-Normalized)')
plt.plot(k_values, accuracies_sklearn, marker='s', label='Sklearn KNN', linestyle='--')
plt.xlabel('k (odd values)')
plt.ylabel('Validation Accuracy')
plt.title('Manual vs Sklearn KNN - Accuracy Comparison')
plt.legend()
plt.grid(True)
plt.show()
```



Final Evaluation on Test Set

We evaluate both models using the test set.

```
print("\n--- Final Evaluation on Test Set ---")
print("Manual KNN Implementation:")
best_k_z_norm=11
print(f"Best k: {best_k_z_norm}")
best_knn_manual = KNeighborsClassifier(n_neighbors=best_k_z_norm)
best_knn_manual.fit(x_z_norm_train, y_z_norm_train)

y_test_pred_manual = best_knn_manual.predict(x_z_norm_test)

test_acc_manual = accuracy_score(y_z_norm_test, y_test_pred_manual)
print(f"Test Accuracy Manual Model: {test_acc_manual:.4f}")
print("\n Classification Report:")
print(classification_report(y_z_norm_test, y_test_pred_manual))

print(" Confusion Matrix:")
print(confusion_matrix(y_z_norm_test, y_test_pred_manual))

print("\n\nSklearn KNN Implementation:")
best_k_sklearn=11
print(f"Best k: {best_k_sklearn}")
best_knn_sklearn = KNeighborsClassifier(n_neighbors=best_k_sklearn)
best_knn_sklearn.fit(x_z_norm_train, y_z_norm_train)

y_test_pred_sklearn = best_knn_sklearn.predict(x_z_norm_test)
```

```

test_acc_sklearn = accuracy_score(y_z_norm_test, y_test_pred_sklearn)
print(f"Test Accuracy Sklearn Model: {test_acc_sklearn:.4f}")
print("\n Classification Report:")
print(classification_report(y_z_norm_test, y_test_pred_sklearn))

print(" Confusion Matrix:")
print(confusion_matrix(y_z_norm_test, y_test_pred_sklearn))

```

```

--- Final Evaluation on Test Set ---
Manual KNN Implementation:
Best k: 11
Test Accuracy Manual Model: 0.8266

```

```

Classification Report:
      precision    recall  f1-score   support

     g         0.78        0.90        0.84        1004
     h         0.89        0.75        0.81        1003

 accuracy                   0.83        2007
 macro avg         0.83        0.83        0.83        2007
weighted avg         0.83        0.83        0.83        2007

```

```

Confusion Matrix:
[[907  97]
 [251 752]]

```

```

Sklearn KNN Implementation:
Best k: 11
Test Accuracy Sklearn Model: 0.8266

```

```

Classification Report:
      precision    recall  f1-score   support

     g         0.78        0.90        0.84        1004
     h         0.89        0.75        0.81        1003

 accuracy                   0.83        2007
 macro avg         0.83        0.83        0.83        2007
weighted avg         0.83        0.83        0.83        2007

```

```

Confusion Matrix:
[[907  97]
 [251 752]]

```

Results Summary

```

summary = pd.DataFrame({
    "Metric": ["Accuracy"],
    "Manual KNN": [test_acc_manual],
    "Sklearn KNN": [test_acc_sklearn]
})
display(summary)

```

Metric	Manual KNN	Sklearn KNN
Training Accuracy	0.826607	0.826607
Testing Accuracy	0.826607	0.826607

This plot compares training and testing accuracy across different k values. It helps visualize: Overfitting (training >> testing for small k)
Underfitting (both low for large k) Optimal k region (balanced accuracies)

```
# Evaluate both models on training and test sets
train_acc_manual = accuracy_score(y_z_norm_train, best_knn_manual.predict(x_z_norm_train))
train_acc_sklearn = accuracy_score(y_z_norm_train, best_knn_sklearn.predict(x_z_norm_train))

# Bar plot comparison
labels = ['Manual KNN', 'Sklearn KNN']
train_acc = [train_acc_manual, train_acc_sklearn]
```